TTIC 31190: Natural Language Processing
Assignment 1: Distributional Word Representations

Instructor: Kevin Gimpel
Assigned: Monday, March 26, 2018
Due: 6:00 pm, Wednesday, April 11, 2018
Submission: email to `kgimpel@ttic.edu`

Submission Instructions

Package your report and code (but not the data) in a single zip file or tarball, name the file with your last name followed by "_hw1", and email the file to `kgimpel@ttic.edu` by 6:00 pm on Wed., April 11.

Collaboration Policy

You are welcome to discuss assignments with others in the course, but solutions and code must be written individually.

Distributional Word Representations

You will implement and experiment with ways of creating word vectors. Let's begin by defining a **vocabulary** V and a **context vocabulary** V_C. The steps below will ask you to implement different ways of building word vectors for the words in V, using context words from V_C. I will provide you a corpus of sentences from English Wikipedia, along with files that you can use to fill V and V_C. Before describing the ways of computing word vectors, we will discuss evaluation and the provided data files.

EVALWS(M): You will quantitatively evaluate your word vectors by comparing the word similarities they produce to human-annotated word similarities. That is, you will be provided with files containing pairs of words along with their human-annotated similarity scores. You should use your word vectors (represented in the matrix M) to compute the cosine similarity of each word pair in the word similarity files. Use Spearman's rank correlation coefficient (also called Spearman's ρ; see `scipy.stats.spearmanr`) as the evaluation metric. That is, for each word pair in a word similarity dataset, you will compute its cosine similarity using your word vectors, which forms a list of similarities. Then, form another list of similarities for those same word pairs, this time containing the manually-annotated similarities from the word similarity datasets we will provide. Compute Spearman's ρ between the two lists of similarities. Do this for two word similarity datasets: MEN and SimLex-999. The files are provided for you. (If any word pairs contain a word that does not appear in the training corpus, simply use a cosine similarity of 0.0 for that pair.)

Provided Data

The following data files are provided to you for this assignment:

- `wiki-{0.1,1}percent.txt` - samples of English Wikipedia downloaded May 15, 2015, with punctuation separated from words, all characters converted to lowercase, and with one sentence per line. For any results and analysis reported in your submission, you should use `1percent`. The smaller file is provided for debugging and preliminary testing. It should still produce reasonable nearest neighbors for most common words.

- `men.txt` - MEN word similarity dataset. The first line is a header, then each line has the format "word1`[tab]`word2`[tab]`similarity_score".

- `simlex-999.txt` - SimLex-999 word similarity dataset. Same format as MEN.

- `vocab-wordsim.txt` - vocabulary file containing the words from the word similarity evaluations (one word per line). Use `vocab-wordsim.txt` for V when you are computing word similarity correlations, but not when you are finding nearest neighbors.

- `vocab-25k.txt` - vocabulary file containing the 25000 most common words in Wikipedia. When computing word similarity correlations, use `vocab-wordsim.txt` for V and `vocab-25k.txt` for V_C. When finding nearest neighbors, use `vocab-25k.txt` for both V and V_C.

1. Required Tasks (50 points total)

1.1 Distributional Counting (20 points)

Implement distributional counting for word vectors. Build a **word-context** matrix C of counts, with size $|V| \times |V_C|$. Each row of C corresponds to a word in V. Each column of C corresponds to a word in V_C. A particular entry C_{ij} in C should equal the number of times that context word j appears within w words of word i in the Wikipedia corpus. A context window contains w words to either side of the center word, so it contains $2w + 1$ words in total (including the center word). Use $w = 3$. Use `vocab-wordsim.txt` to populate V and use `vocab-25k.txt` to populate V_C. When computing counts, be sure to pad the sentences with the start-of-sentence symbol `<s>` and end-of-sentence symbol `</s>`. (These are both included as words in the beginning of the `vocab-25k.txt` file.) Call EVALWS(C) and report your results on both MEN and SimLex-999. As a sanity check, your Spearman correlation for MEN should be close to 0.22. Submit your code.

1.2 Computing PMIs (10 points)

After building C as above, build a new matrix C_{pmi} that contains pointwise mutual information (PMI) values for each entry (i.e., each word pair). See Section 4.2 of Turney and Pantel (2010) for more details. Call EVALWS(C_{pmi}) and report your results. You should see improvements in Spearman correlation for both MEN and SimLex-999. Submit your code.

1.3 Experimentation (5 points)

Report the results of EVALWS(C) and EVALWS(C_{pmi}) for three values of w: $1, 3$, and 6. You should notice systematic trends in terms of correlation as window size changes. Look at the manually-annotated similarities in the MEN and SimLex-999 datasets and describe why you think the two datasets show the trends they do.

1.4 Analysis (15 points)

In this component, you will analyze your word vectors qualitatively by computing nearest neighbors for particular query words and comparing them across window sizes. You may also find it helpful to manually inspect instances of certain words in the original corpus to find typical contexts for each word.

For this section, use `vocab-25k.txt` to populate both V and V_C. Compute the PMI matrix C_{pmi} with window sizes $w = 1$ and $w = 6$. To find the n nearest neighbors for a query word, compute cosine similarity between the query word and all words in V and then keep the n words with highest cosine similarity. When doing so, omit the query word from the list of nearest neighbors since it will always have cosine similarity of 1.

1.4.1 Warm-up: Printing nearest neighbors (3 points)

For the two window sizes $w = 1$ and $w = 6$, compute and print the 10 nearest neighbors for the query word *monster*. (Hint: using my implementation, the nearest neighbor of *monster* is *dragon* with $w = 1$ and *evil* with $w = 6$. Your nearest neighbor lists may differ slightly from mine, but hopefully these two are high up in your lists.)

1.4.2 Part-of-speech tag similarity in nearest neighbors (7 points)

Do nearest neighbors tend to have the same part-of-speech tag as the query word, or do they differ? Does the pattern differ across different part-of-speech tags for the query word? How does window size affect this?

Explore these questions by choosing query words with different parts of speech and computing their nearest neighbors. When choosing query words, consider nouns, verbs, adjectives, and prepositions. (Hint: when considering verbs, use inflected forms like *transported*.) Try a few query words from each part of speech category and see if you can find any systematic patterns when comparing their nearest neighbors across window sizes $w = 1$ and 6. When the neighbors differ between window sizes, how do they differ? Can you find any query words that have almost exactly the same nearest neighbors with the two window sizes?

Discuss your findings, showing examples of nearest neighbors for particular words to support your claims.

1.4.3 Words with multiple senses (5 points)

Now try choosing words with multiple senses (e.g., *bank*, *cell*, *apple*, *apples*, *axes*, *frame*, *light*, *well*, etc.) as query words. What appears to be happening with these words? What happens when you compare the neighbors with different window sizes ($w = 1$ vs. $w = 6$)? Discuss your findings, showing examples of nearest neighbors for particular words to support your claims.

2. Extra Credit: Dimensionality Reduction (5 extra points possible)

The following additional data files are provided to you for the extra credit:

- `vocab-25k+wordsim.txt` - vocabulary file containing the top 25000 most common words in Wikipedia as well as the words contained in the word similarity evaluations. Use this file for V for this extra credit component only.

- `vocab-3k.txt` - vocabulary file containing the 3000 most common words in Wikipedia. Use this file for V_C for this extra credit component only.

Compute the PMI matrix C_{pmi} and then compute the **truncated SVD** of C_{pmi}. That is, compute the singular value decomposition (SVD) of C_{pmi}, then retain only the top k dimensions. See Section 19.5 of *Speech and Language Processing (3rd Edition)* for more details. Most programming languages will have free libraries that implement SVD.[1]

Use $w = 3$ and experiment with $k \in \{10, 50, 100, 500, 1000\}$. EVALWS the resulting matrix for each k and report your results in a table for both MEN and SimLex-999 and all k values. Note that the expense of performing the SVD will depend on the sizes of V and V_C, so make sure that you use the given vocabulary files to prune these vocabularies. In particular, use `vocab-25k+wordsim.txt` to populate V and `vocab-3k.txt` to populate V_C.

Since you will be using a different context vocabulary from the evaluations in Sections 1.2-1.3 above, you should compute new baseline Spearman correlations using the V and V_C defined in this section to serve as a proper baseline to compare to your truncated SVD results.

3. Implementation Tips

The most difficult part of this assignment is making your implementation efficient enough to scale up to large corpora. Below are some tips to keep in mind:

- You only need to do one pass through the corpus to compute C (or C_{pmi}), so you do not need to store the corpus in memory.

- Though the notation above talked about representing the word/context counts in one big matrix C (or C_{pmi}), this is probably not feasible in practice. You will likely need to use sparse data structures to store the counts (these may have various names depending on the programming language, e.g., maps, hashes, or dictionaries). For example, you can use a sparse data structure to map a word (from V) to a sparse data structure that maps context words (from V_C) to counts. While this may not be strictly necessary when computing word similarity correlations (since you can use the small file `vocab-wordsim.txt` for V), this will be necessary when finding nearest neighbors since you cannot restrict the vocabulary as much. You should also exploit sparsity when computing cosine similarities.

- If you encounter underflow when estimating the small probabilities used for PMI calculation, you can perform the computation in the log domain (though I did not have to do this for my implementation).

References

Turney, P. D. and Pantel, P. (2010). From frequency to meaning: Vector space models of semantics. *J. Artif. Int. Res.*, 37(1):141–188. [2]

---

[1]In my solution, I used C++ and Eigen (http://eigen.tuxfamily.org/index.php?title=Main_Page).